

Support ticket dataset successfully cleaned and ready for classifier training, with quality score improved from 71.2 to 98.7 out of 100.

The support tickets dataset was thoroughly processed, reducing the document count from 148 to 134 while resolving critical issues like PII leaks and cross-label duplicates. Sensitive personal data was found in 18 documents, and 57 documents contained a repeated confidentiality notice that was successfully stripped. Although cleaning drastically improved overall quality, 4 documents remain flagged as likely mislabels due to centroid mismatches. The dataset is now well-prepared for training the support ticket classifier.

Data quality before	After cleaning	Documents	Numbers fact-checked
71.2/100	98.7/100	148 → 134	17

What this dataset is: This is a corpus of 148 customer support text documents categorized by folder into three labels (billing, shipping, and technical) intended for training a ticket classification model.

Findings

WARNING Confidentiality boilerplate appears across documents

A repeated confidentiality notice was present in 57 documents, representing 38.51% of all raw documents.

Recommendation: Strip boilerplate disclaimers prior to feature extraction so vocabulary and classifier weights reflect actual ticket issues.

Evidence: txt_boilerplate_001

CRITICAL Personal data detected across ticket documents

A total of 18 documents contained sensitive personal information, including 10 with phone numbers, 6 with email addresses, and 2 with credit card numbers.

Recommendation: Mask or redact all personal data tokens before training to protect user privacy and avoid model memorization.

Evidence: txt_pii_001

CRITICAL Duplicate tickets identified including cross-label group

The raw data contained 5 near-duplicate extra copies, including 1 duplicate group that spanned across distinct category labels.

Recommendation: Remove duplicate documents and resolve cross-label conflicts to prevent conflicting supervision during model training.

Evidence: txt_dupes_001

WARNING Several documents read closer to other labels

A TF-IDF centroid check showed 97.12% centroid agreement, while 4 documents were flagged as likely mislabels.

Recommendation: Manually review the 4 flagged documents and relabel them if confirmed misclassified.

Evidence: txt_labels_001

WARNING Tickets contain foreign languages and encoding damage

Raw data contained 5 documents in other languages and 3 documents with character encoding damage.

Recommendation: Filter out non-English tickets for a monolingual English classifier and repair character encodings.

Evidence: txt_quality_001

INFO Data cleaning improved quality score significantly

Cleaning reduced the document count from 148 to 134, removing 14 documents and increasing the quality score from 71.2 to 98.7 out of 100.

Recommendation: Proceed with the cleaned dataset of 134 tickets for downstream text classification modeling.

Evidence: cleaning_summary_001

Next steps

1. Review the 4 documents flagged as likely mislabels before finalizing the training split.
2. Verify the masking of sensitive personal data across the 18 affected documents.
3. Proceed with training the support ticket classifier using the cleaned dataset of 134 documents.

How the agents checked themselves

Model calls 6 · guardrail catches 0 · self-corrections 0 · review rounds 1 · model switches 4

+61s Switched models for the reviewer

gemini-3.8-flash hit its server error 503 (set aside for 30s); trying the next model in the pool.

+65s Switched models for the reviewer

gemini-3.7-flash hit its server error 503 (set aside for 30s); trying the next model in the pool.

+70s Switched models for the reviewer

gemini-3.6-flash hit its server error 503 (set aside for 30s); trying the next model in the pool.

+72s Switched models for the reviewer

gemini-3.5-flash hit its server error 503 (set aside for 30s); trying the next model in the pool.

+74s Reviewer approved the findings

Cleaning applied

Repair encoding and HTML markup, normalize whitespace, strip repetitive confidentiality boilerplate, and mask sensitive PII. Deduplicate texts across labels, remove short documents, filter out non-English tickets, and flag suspected misclassifications for human review.


#	Operation	Risk	Columns	Documents	Why
1	fix_encoding	safe	all	148	Repair character encoding damage found in corrupted documents to restore text readability.
2	strip_html	safe	all	148	Remove HTML markup and decode entities found across ticket documents.
3	normalize_whitespace	safe	all	148	Collapse redundant whitespace and empty lines introduced during formatting and tag stripping.
4	strip_boilerplate	lossy	all	148	Remove widespread confidentiality notice boilerplate occurring in over 38% of documents.
5	mask_pii	lossy	all	148	Mask sensitive personal data including emails, phone numbers, card numbers, and SSNs in 18 documents.
6	drop_exact_duplicates	destructive	all	148 → 134	Remove identical redundant copies across 4 exact duplicate groups.

#	Operation	Risk	Columns	Documents	Why
7	drop_near_duplicates	destructive	all	143 → 142	Deduplicate near-identical documents exceeding the 0.8 similarity threshold.
8	drop_cross_label_duplicates	destructive	all	142	Eliminate conflicting duplicate documents that appear under multiple distinct category labels.
9	drop_short_documents	destructive	all	142 → 139	Drop empty documents and low-information tickets containing fewer than 5 words.
10	filter_language	destructive	all	139 → 134	Keep primary English documents while filtering out detected foreign language tickets.
11	flag_suspected_mislabels	safe	all	134	Flag four documents identified by centroid mismatch as reading closer to another class for review.


The same steps are in `cleaning_pipeline.py`, which reruns on the full data.

Charts


Documents per label




Document length




Most frequent words



Languages



Documents per label



Label balance

Class	Documents	Share %
billing	57	38.5135
shipping	52	35.1351
technical	39	26.3514

Largest class / smallest class: 1.4615

Documents that read like another label (checked by code)

Document	Label	Reads like	Text (masked)
billing/ticket_bill_052.txt	billing	shipping	Hello, My package has not arrived and the tracking number says delivered. The courier left the pa...
technical/ticket_tech_032.txt	technical	billing	Hello team, I was charged twice for my Team Workspace subscription this month. Why is there a lat...

Document	Label	Reads like	Text (masked)
technical/ticket_tech_033.txt	technical	billing	Hello team, I was charged twice for my subscription and need a refund of the extra payment on my...
shipping/ticket_ship_047.txt	shipping	technical	Hi, The app crashes every time I log in after the latest update, and sync stopped working with an...

Topics

Topic	Strongest words	Share %
1	payment, card, email, charge, billing, notice	38.5135
2	delivery, parcel, order, left, tracking, address	33.1081
3	returns, timeout, work, stopped, files, api	25.0

Every number in the findings was checked against the evidence artifacts listed. Models used: gemini-3.5-flash-lite, gemini-3.8-flash.